

Q. Reconfigure the existing deployment front-end and add a port specification named http , exposing port 80/tcp of the existing container nginx.

Q. Create a New service named front-end-svc exposing the container port http

Q. Configure new service to also expose the individual pods via a NodePort on the nodes on which they are scheduled.

=====

Q. Create a pod with nginx image, named as web in dev-test namespace
SOLU

=====

```
$ kubectl create ns dev-test
$ kubectl run web --image=nginx --restart=Never -n dev-test
```

=====

Q. show the pods that are part of namespace XYZ , show services that is part of namespace ABC

SOLU

=====

```
$ kubectl get po -n XYZ
$ kubectl get svc -n ABC
```

=====

```
kube3a k top po -n kube-system | head -3
NAME                                CPU(cores)   MEMORY(bytes)
coredns-86c58d9df4-562mf           1m           7Mi
coredns-86c58d9df4-zl4mj           1m           7Mi
```

List the top 1 pod with the highest CPU usage

```
kubectl top po -n kube-system | sort -nr -k2 | head -1
^^<== -k is column number , so
column 2 is the CPU
```

List the top 1 pod with the highest MEM usage

```
kubectl top po -n kube-system | sort -nr -k3 | head -1
^^<== -k is column , Column 3 is
the Memory
```

=====

1. Create a static pod on a node and put the pod specification file in /etc/kubernetes/manifests.

SOLU

=====

Refer to <https://kubernetes.io/docs/tasks/administer-cluster/static-pod/> , and use the static-web.yaml

```
root@kube3b:/etc/kubernetes/manifests# pwd
/etc/kubernetes/manifests
root@kube3b:/etc/kubernetes/manifests# cat static-web.yaml
apiVersion: v1
kind: Pod
metadata:
  name: static-web
  labels:
    role: myrole
spec:
  containers:
    - name: web
      image: nginx
      ports:
        - name: web
          containerPort: 80
          protocol: TCP
root@kube3b:/etc/kubernetes/manifests#
```

Refer to : <https://medium.com/@sonasingh46/static-pod-in-kubernetes-e3854507655f>
vi /etc/systemd/system/kubelet.service

[Service]

```
ExecStart=/home/kubernetes/bin/kubelet $KUBELET_OPTS \
    --pod-manifest-path=/etc/kubernetes/manifests
```

```
# systemctl daemon-reload ; kubect1 get po ; if no pod then
# systemctl restart kubelet
```

=====

2. Create a deployment nginx-dns and expose it as service nginx-dns.

Create busybox pod , Use nslookup to get dns result for service & pod and store the values into a file.

SOLU

=====

```
kube3a k run nginx-dns --image=nginx --port=80
```

```
deployment.apps/nginx-dns created
kube3a k expose deployment nginx-dns --port=80 --target-port=80
service/nginx-dns exposed
kube3a
```

```
by default nginx assings label as run=deploymet|pod_name>
```

```
kube3a k get deploy,po,svc -l run=nginx-dns -o wide
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
CONTAINERS IMAGES SELECTOR				
deployment.extensions/nginx-dns	1/1	1	1	4m44s
nginx run=nginx-dns				

NAME	READY	STATUS	RESTARTS	AGE	IP
NODE NOMINATED NODE READINESS GATES					
pod/nginx-dns-6b7f858b9b-hzcv	1/1	Running	0	4m44s	10.244.2.153
kube3c.home.local	<none>	<none>			

```
^^^^^^^^^^<== we need to resolve this for pod.dns
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
SELECTOR					
service/nginx-dns	ClusterIP	10.98.27.246	<none>	80/TCP	3m47s
run=nginx-dns					

```
kube3a ^^^^^^^^^^^<== we need to resolve this for
service.dns
```

```
kube3a k run bbox --image=busybox:1.28 --restart=Never -- /bin/sh -c 'sleep 360000'
pod/bbox created
```

```
Resolve pod IP :
kube3a k exec bbox -- nslookup 10.244.2.153 > /tmp/aa
kube3a
Edit /tmp/aa and delete everything except :
10-244-2-153.nginx-dns.default.svc.cluster.local
```

```
kube3a sudo cp /tmp/aa /opt/TTT/pod.dns
kube3a cat /opt/TTT/pod.dns
10-244-2-153.nginx-dns.default.svc.cluster.local
```

```
Next to resolve service for pod :
kube3a k exec bbox -- nslookup nginx-dns > /tmp/bb
kube3a
```

```
Edit /tmp/bb and ensure you ONLY have :
kube3a cat /tmp/bb
      nginx-dns.default.svc.cluster.local
kube3a
```

```
kube3a sudo cp /tmp/bb /opt/TTT/service.dns
kube3a cat /opt/TTT/service.dns
      nginx-dns.default.svc.cluster.local
kube3a
```

=====

3. Create a init container which shouldn't start if a file is not found.
SOLU
=====

```
I created :
$ k run bbox --image=busybox --restart=Never --dry-run -o yaml -- /bin/sh -c 'if [
! -f /opt/XX/xx]
>then
>exit 0
>else
>sleep 3600000000
>fi' > bbox12.yaml
```

Once I had above then I added entries for volume: and volumeMounts: for emptyDir,
then

```
dharm@kube3a:~$ cat bbox12.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: bbox
  name: bbox
spec:
  volumes:
    <== U need emptyDir: or hostpath:
  - name: cache-volume
    emptyDir: {}
  containers:
  - args:
    - /bin/sh
    - -c
    - |-
```

```

    if [ ! -f /opt/XX/xx ]
    then
    exit 0
    else
    sleep 36000000
    fi
image: busybox
name: bbox
volumeMounts:
- mountPath: /opt/XX <== the initcontainer and the actual container shares
same dir
  name: cache-volume
initContainers:
- args:
  - /bin/sh
  - -c
  - 'touch /opt/XX/xx'
image: busybox
name: bbox1
volumeMounts:
- mountPath: /opt/XX <== the initcontainer and the actual container shares
same dir
  name: cache-volume

#
## dnsPolicy: ClusterFirst
## restartPolicy: Never

```

```

kube3a k create -f bbox12.yaml
pod/bbox created

```

```

kube3a k get po --watch

```

NAME	READY	STATUS	RESTARTS	AGE
bbox	0/1	Init:0/1	0	12s
bbox	0/1	PodInitializing	0	16s
bbox	1/1	Running	0	18s
bbox	0/1	Completed	0	19s
bbox	1/1	Running	1	21s
bbox	0/1	Completed	1	22s

^

```

kube3a k logs bbox
FILE PRESENT BOSSS

```

=====

4. Bootstrap a kubernetes node.

<https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-tls-bootstrapping/>

=====

5. Use node selector to assign a pod to particular node.

SOLU

=====

Label a Node :

```
$ kubectl label no kube3c    node=kube3c
```

Create alpine pod , and assign to a Node labeled node=kube3c :

```
kube3a k run alpine --image=alpine --restart=Never --dry-run -o yaml --command sleep 36000000 > alpine-nodeselector.yaml
```

```
kube3a cat alpine-nodeselector.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  creationTimestamp: null
```

```
  labels:
```

```
    run: alpine
```

```
  name: alpine
```

```
spec:
```

```
  containers:
```

```
  - command:
```

```
    - sleep
```

```
    - "36000000"
```

```
    image: alpine
```

```
    name: alpine
```

```
    resources: {}
```

```
  dnsPolicy: ClusterFirst
```

```
  restartPolicy: Never
```

```
  nodeSelector:          <== edit file and add this manually
```

```
    node: kube3c         <== add this and add this manually
```

```
kube3a
```

```
kube3a k create -f alpine-nodeselector.yaml
```

```
pod/alpine created
```

```
kube3a k get po alpine -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED	NODE	READINESS GATES				

```
alpine      0/1      ContainerCreating    0           13s    <none>    kube3c.home.local
<none>      <none>
kube3a
```

=====

6. Label a node and it shouldn't be allowed to execute/schedule any pods.

SOLU

=====

At present :

Taints: node-role.kubernetes.io/master:NoSchedule

Taints: <none>

Taints: <none>

kube3a

kube3a k taint node kube3c aa=bb:NoSchedule

kube3a k describe no | egrep -i taint

Taints: node-role.kubernetes.io/master:NoSchedule

Taints: <none>

Taints: aa=bb:NoSchedule

kube3a

kube3a k get po --all-namespaces -o wide | egrep kube3c | wc -l

2

Next deploy 20 nginx pods and see if any gets deployed on kube3c :

kube3a k get po --all-namespaces -o wide | egrep -i kube3c | wc -l

2

AS per above we still have 2 pods that are running on kube3c , as those are part of daemonset

=====

7. Create a deployment name nginx-app with 3 replicas use nginx:1.9.1, do a rolling update to nginx:1.13.5 and then undo it.

New image was asked to use for rolling update.

SOLU

=====

kube3a k run nginx-app --image=nginx:1.9.1 --replicas=3

deployment.apps/nginx-app created

kube3a k rollout status deploy nginx-app

deployment "nginx-app" successfully rolled out

```
kube3a k describe deploy nginx-app | egrep -i image:
      Image:          nginx:1.9.1
kube3a
```

Next update the nginx to ver 1.13.5

```
kube3a k set image deployment nginx-app  nginx-app=nginx:1.13.5
deployment.extensions/nginx-app image updated
```

```
kube3a k rollout status deploy nginx-app
```

Waiting for deployment "nginx-app" rollout to finish: 1 out of 3 new replicas have been updated...

Waiting for deployment "nginx-app" rollout to finish: 1 out of 3 new replicas have been updated...

Waiting for deployment "nginx-app" rollout to finish: 1 out of 3 new replicas have been updated...

Waiting for deployment "nginx-app" rollout to finish: 2 out of 3 new replicas have been updated...

Waiting for deployment "nginx-app" rollout to finish: 2 out of 3 new replicas have been updated...

Waiting for deployment "nginx-app" rollout to finish: 2 out of 3 new replicas have been updated...

Waiting for deployment "nginx-app" rollout to finish: 1 old replicas are pending termination...

Waiting for deployment "nginx-app" rollout to finish: 1 old replicas are pending termination...

deployment "nginx-app" successfully rolled out
kube3a

```
kube3a k describe deployment nginx-app | egrep -i image:
      Image:          nginx:1.13.5
kube3a
```

Undo the upgraqde to version 1 (previous) :

```
kube3a k rollout  undo deployment nginx-app --to-revision=1
deployment.extensions/nginx-app rolled back
```

```
kube3a k describe deploy nginx-app | egrep -i image:
      Image:          nginx:1.9.1
kube3a
```

=====

8. Node is in not ready status. Find the issue and fix it.

SOLU

=====

ssh onto the worker Node , and run :

```
$ sudo systemctl status -l kubelet
```

If the kubelete service is not runing then :

```
$ sudo systemctl enable kubelet
```

```
$ sudo systemctl restart kubelet
```

Finally run below to ensure kubelet service is runing fine :

```
$ sudo systemctl status -l kubelet
```

=====

9. Find all pods which are using same service in a namespace and store the value into a file.

SOLU

=====

Create a ns called boss :

```
kube3a k create ns boss
```

```
namespace/boss created
```

```
kube3a k get ns boss
```

```
NAME      STATUS    AGE
```

```
boss      Active    10s
```

```
kube3a
```

deploy busybox pods and expose port 80 and create service in ns boss.

```
kube3a k get po -n boss
```

```
NAME      READY    STATUS    RESTARTS   AGE
```

```
b         1/1      Running   0           2m27s
```

```
b1        1/1      Running   0           2m22s
```

```
b2        1/1      Running   0           2m15s
```

```
b3        1/1      Running   0           109s
```

```
b4        1/1      Running   0           103s
```

```
b5        1/1      Running   0           98s
```

```
b6        1/1      Running   0           86s
```

```
kube3a
```

above is total 8 , but then we ignore the header "NAME READY STATUS RESTARTS AGE" ,
si its 7

```
$ sudo echo 7 > /opt/file/X
```

```

kube3a k get svc -n boss
NAME      TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
b3        ClusterIP     10.96.172.82    <none>           80/TCP           8m
b4        ClusterIP     10.107.213.233  <none>           80/TCP           7m54s
b5        ClusterIP     10.98.80.82     <none>           80/TCP           7m49s
b6        ClusterIP     10.102.82.0     <none>           80/TCP           7m37s
kube3a k get svc -n boss | wc -l
5

```

By ignoring headers , we have 4 :

```
$ sudo echo 4 > /opt/file/XX
```

10. Find all pods with label cpu=loader and sort them by consuption and store the value into a file.

SOLU

====

I have create couple of pods , all have label cpu=loader, lets see which has High CPU Usage
and which consumes more Memory.,

```

kube3a k top po -l cpu=loader
NAME      CPU(cores)    MEMORY(bytes)
a1        0m            2Mi
a2        1m            4Mi
a3        6m            14Mi
a4        1m            1Mi
a5        0m            0Mi
a6        0m            0Mi
kube3a

```

```

echo "a3" > high.cpu
echo "a3" > high.mem

```

Delete all the above pods that have label cpu: loader

```

kube3a k delete po -l cpu=loader
pod "a1" deleted
pod "a2" deleted
pod "a3" deleted
pod "a4" deleted
pod "a5" deleted
pod "a6" deleted

```

kube3a

=====

11. Find all persistent volumes in a namespace and sort them by capacity and store the value into a file.

SOLU
=====

```
kube3a k get pv --sort-by=.spec.capacity.storage
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM
STORAGECLASS  REASON    AGE
pv-1g         1Gi       RWO           Retain          Available
103s
myvolume      10Gi      RWO,RWX       Retain          Available  normal
4m5s
myvolume1     15Gi      RWO,RWX       Retain          Available  normal
3m11s
pv-1f         100Gi     RWO           Retain          Available
57s
kube3a
```

=====

12. Create a pod with multiple containers including image of redis, nginx , memcached, consul, busybox and alpine.

SOLU
=====

```
NOTE : ONLY busybox and alpine requires "--command sleep 3600000" , and rest does NOT
$ k run pod-all --image=alpine --restart=Never --dry-run -o yaml --command sleep 3600000 > pod-all.yaml
```

```
kube3a cat pod-all.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod-all
  name: pod-all
```

```
spec:
  containers:
  - command:
    - sleep
    - "3600000"
    image: alpine
    name: p1
#
  - command:
    - sleep
    - "3600000"
    image: busybox
    name: p2
#
  - image: redis
    name: p3
#
  - image: memcached
    name: p4
#
  - image: nginx
    name: p5
#
  - image: consul
    name: p6
#
  dnsPolicy: ClusterFirst
  restartPolicy: Never
#
kube3a
kube3a k create -f pod-all.yaml
    pod/pod-all created
```

After few mins..

```
kube3a k get po
NAME      READY   STATUS    RESTARTS   AGE
pod-all   6/6     Running   0           4m52s
kube3a
```

=====

13. Create a secret super-secret with credential as bob and create two pods, one which mount secret as volume and another which mount it as env variable.

SOLU

=====

```
kube3a k create secret generic super-secret --from-literal=credential=bob
secret/super-secret created
```

```
kube3a k get secret super-secret
```

NAME	TYPE	DATA	AGE
super-secret	Opaque	1	76s

```
kube3a k get secret super-secret --export -o yaml
```

```
apiVersion: v1
```

```
data:
```

```
  credential: Ym9i      <== echo Ym9i | base64 -d , will give you bob
```

```
kind: Secret
```

```
metadata:
```

```
  creationTimestamp: null
```

```
  name: super-secret
```

```
  selfLink: /api/v1/namespaces/default/secrets/super-secret
```

```
type: Opaque
```

```
kube3a
```

Create environment variable called TOPSECRET and assign the secret value , refer to :

<https://kubernetes.io/docs/concepts/configuration/secret/>

```
$ pod-secret.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: secret-env-pod
```

```
spec:
```

```
  containers:
```

```
    - name: mycontainer
```

```
      image: redis
```

```
      env:
```

```
        - name: TOPSECRET
```

```
          valueFrom:
```

```
            secretKeyRef:
```

```
              name: super-secret
```

```
              key: credential
```

In the above secret url , I searched for env: and saw above , which I cp & paste

```
kube3a k create -f pod-secret.yaml
```

```
pod/secret-env-pod created
```

kube3a

kube3a k get po

NAME	READY	STATUS	RESTARTS	AGE
secret-env-pod	1/1	Running	0	80s

kube3a

kube3a k exec secret-env-pod -- printenv | egrep TOPSECRET
TOPSECRET=bob
kube3a

Next we need to copy the secret in a volume :

kube3a cat secret-pod.yaml

```
kind: Pod
apiVersion: v1
metadata:
  name: secret-pod
  labels:
    name: secret-pod
spec:
  volumes:
  - name: secret-volume
    secret:
      secretName: super-secret
  containers:
  - name: secret-pod
    image: nginx
    volumeMounts:
    - name: secret-volume
      readOnly: true
      mountPath: "/etc/secret-volume"
```

kube3a

From the secret url , I searched for volume: , and found above , which I did minor changes and bingo..

kube3a k create -f secret-pod.yaml
pod/secret-pod created

kube3a

kube3a k exec secret-pod -- ls -l /etc/secret-volume
lrwxrwxrwx 1 root root 17 Mar 19 22:58 credential -> ../data/credential

kube3a k exec secret-pod -- cat /etc/secret-volume/credential
bob

=====

14. Fix a malfunctioning kubernetes cluster.

=====

15. Create a persistent volume with 1GB details provided.

SOLU

=====

On k8s.io , seach for "persistentvolume as hostpath" , will see below :

\$ pv-1G.yaml

kind: PersistentVolume

apiVersion: v1

metadata:

name: task-pv-volume

labels:

type: local

spec:

storageClassName: manual

capacity:

storage: 1Gi

accessModes:

- ReadWriteOnce

hostPath:

path: "/mnt/data"

kube3a k create -f pv-1G.yaml

persistentvolume/task-pv-volume created

kube3a k get pv

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
task-pv-volume	1Gi	RWO	Retain	Available	
manual		2s			

kube3a

=====

Q 15.1 - Create a pv of size 1GB in namespace qa , using hostpath /srv/boss

SOLU

=====

=====

16. Create a pod nginx and expose it.

SOLU

=====

```
$ k run nginx --image=nginx --restart=Never --port=80 --expose
```

Above will create nginx pod , with port 80 open in firewall , and will expose as ClusterIP

=====

17. Find all nodes in a cluster which have no taints and store the value into a file.

SOLU

=====

```
kube3a k describe no | egrep -i taint
Taints:                node-role.kubernetes.io/master:NoSchedule
Taints:                <none>
Taints:                <none>
kube3a
```

```
$ sudo echo "2" > /opt/XX/k34ft.txt
```

=====

18. Create a namespace website-backend, run jenkins pod inside it and expose it.

SOLU

=====

```
$ k create ns website-backend
$ k run jenkins --image=jenkins --restart=Never -n website-jenkins --port=80
--expose
```

=====

19. ETCD Backup

<https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>

SOLU

=====

```
kube3a pwd
/home/dharmin/etcd-v3.3.12-linux-amd64
```



```
kube3a ./etcdctl -v
etcdctl version: 3.3.12
API version: 2
kube3a
```

By default a etcd is in kube-system namespace

```
kube3a k get po --all-namespaces | awk '$2 ~ /etcd/ {print $2}'
      etcd-kube3a
kube3a
```

```
kube3a k get po etcd-kube3a -n kube-system --export -o yaml | egrep "\-file|data-
dir|endpoint"
- --cert-file=/etc/kubernetes/pki/etcd/server.crt          <== NEED THIS
- --data-dir=/var/lib/etcd                                <== NEED THIS
- --key-file=/etc/kubernetes/pki/etcd/server.key           <== NEED THIS
- --peer-cert-file=/etc/kubernetes/pki/etcd/peer.crt
- --peer-key-file=/etc/kubernetes/pki/etcd/peer.key
- --peer-trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt
- --trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt       <== NEED THIS
kube3a
```

```
kube3a ./etcdctl backup --help
```

NAME:

etcdctl backup - backup an etcd directory

USAGE:

etcdctl backup [command options]

OPTIONS:

--data-dir value	Path to the etcd data dir	<== NEED THIS
--wal-dir value	Path to the etcd wal dir	
--backup-dir value	Path to the backup dir	<== NEED THIS
--backup-wal-dir value	Path to the backup wal dir	
--with-v3	Backup v3 backend data	

```
kube3a
```

```
kube3a sudo ./etcdctl --cert-file /etc/kubernetes/pki/etcd/server.crt --key-file
/etc/kubernetes/pki/etcd/server.key \
      --ca-file /etc/kubernetes/pki/etcd/ca.crt backup --data-dir
/var/lib/etcd --backup-dir /tmp/ABC
[sudo] password for dharmin:
2019-03-20 00:05:36.754230 I | ignoring v3 raft entry
2019-03-20 00:05:36.754479 I | ignoring v3 raft entry
```

```
2019-03-20 00:05:36.754856 I | ignoring v3 raft entry
2019-03-20 00:05:36.755113 I | ignoring v3 raft entry
```

Lets check out backed-up data :

```
kube3a sudo chmod -R 755 /tmp/ABC
kube3a
```

```
kube3a find /tmp/ABC -print
/tmp/ABC
/tmp/ABC/member
/tmp/ABC/member/wal
/tmp/ABC/member/wal/00000000000000000-00000000000000000.wal
/tmp/ABC/member/snap
/tmp/ABC/member/snap/00000000000000007-000000000000a3973.snap
/tmp/ABC/member/snap/db
kube3a
```

Lets check if the backed up data is good :

On k8s.io , search for ETCDCTL_API=v3 , and in the pg you see scroll down and should see :

```
kube3a ETCDCTL_API=3 ./etcdctl --write-out=table snapshot status
/tmp/ABC/member/snap/db
```

HASH	REVISION	TOTAL KEYS	TOTAL SIZE
7fb25e0b	0	1	25 kB

```
kube3a
```

20. Create nginx as daemonset.

SOLU

=====

Remember a DaemonSet is similar to a Deployment, so create a Deployment manifest and change is to a daemonset :

```
kube3a k run redis-ds --image=redis --restart=Always --dry-run -o yaml > redis-ds.yaml
```

^^^^^^^^^^^^^^^^<== this is optional, don't

```
require

$ vi redis-ds.yaml
apiVersion: apps/v1
kind: DaemonSet    <== this needs to be changed from Deployment to DaemonSet
metadata:
  creationTimestamp: null
  labels:
    run: redis-ds
  name: redis-ds
spec:
  replicas: 1    <===== Delete this
  selector:
    matchLabels:
      run: redis-ds
  strategy: {}    <===== Delee this
  template:
    metadata:
      creationTimestamp: null
      labels:
        run: redis-ds
    spec:
      containers:
      - image: redis
        name: redis-ds
        resources: {}    <== delete this
status: {}    <<== Delete this

kube3a k create -f  redis-ds.yaml
      daemonset.apps/redis-ds created

kube3a k get ds,po  -o wide

NAME                                DESIRED    CURRENT    READY    UP-TO-DATE    AVAILABLE
NODE SELECTOR    AGE    CONTAINERS    IMAGES    SELECTOR
daemonset.extensions/redis-ds    2            2            2            2            2
<none>            35s    redis-ds    redis    run=redis-ds

NAME                                READY    STATUS    RESTARTS    AGE    IP                                NODE
NOMINATED NODE    READINESS GATES
pod/redis-ds-45dd5    1/1    Running    0            35s    10.244.2.152
kube3c.home.local    <none>    <none>
pod/redis-ds-tf5qp    1/1    Running    0            35s    10.244.1.164
kube3b.home.local    <none>    <none>
kube3a
```

=====

21. Scale number of pods of a deployment to 6.

SOLU

=====

```
$ k run nginx --image=nginx --replicas=2
$ k scale deployment nginx --replicas=6
```

=====

22. Store the logs of a container having error file-not-found.

SOLU

=====

```
$ k log <pod_name> | egrep "file-not-found" > log1
$ sudo cp log1 /opt/XXX/kdrsv
```

=====

23. Create a pod with local volume details provided.

SOLU

=====

refer to k8s.io and search for hostPath or emptyDir

=====

24. Create a pod and assign it to node labelled as disk=spinning.

SOLU

=====

check if the nodes have the label disk=spinning :

```
$ k get no ---show-labels
```

```
$ k run redis --image=redis --restart=Never --dry-run -o yaml > redis-pod.yaml
then at the end of the line add :
```

```
  nodeSelector:
```

```
    disk: spinning
```

Ensure the "nodeSelector:" is in line with "containers:"

=====

if use busybox or alpine pod then always do use --command sleep 3600000 or --
/bin/sh -c 'command'

when comes to nginx or redis pod's , does not require --command or -- /bin/sh ,
as per above

PRE-REQ
=====

```
$ vi a
alias k=kubectl
alias c=clear

$ source ~/a
$ cp a /tmp
```

When you become root , then :
source /tmp/a

0. Create following deployments with 2 pods : nginx, redis, memcached, consul
expose nginx to service boss
Forget above question

Can you tell which pods are part of service boss ?
SOLU
===

If you create a service from a deployment/pod , the default label for service is
same as of deployment/pod.

Below will show you Selector , which is same as label :
\$ kubectl get svc boss --expose -o yaml

Once you know what label is assinged to above service , then
\$ kubectl get po -l <label_name>

You'll see the pods that use the service boss.

=====

1. Create nginx static pod , copy the yaml file in /etc/kubernetes/manifest dir.
SOLU
=====

On worker Node , as user root :
source /tmp/a
cd /etc/kubernetes/manifests
copy the static pod here from <https://kubernetes.io/docs/tasks/administer-cluster/static-pod/>

```
# systemctl status -l kubelet
kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset:
enabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
           10-kubeadm.conf
   Active: active (running) since Mon 2019-03-11 11:47:55 GMT; 6s ago
     Docs: https://kubernetes.io/docs/home/
  Main PID: 1624 (kubelet)
```

Our worker Node kube3c is online (Ready) :

```
kube3a k get no
```

NAME	STATUS	ROLES	AGE	VERSION
kube3a	Ready	master	5d2h	v1.13.0
kube3b	Ready	<none>	5d1h	v1.13.0
kube3c	Ready	<none>	5d1h	v1.13.0

```
kube3a
```

14. If you need to know what options to use with securityContext in a pod , how can you find out

SOLU

====

```
$ kubectl explain pod.spec.securityContext --recursive
```

15. Assign a Worker Node with taint tz=tanzania:NoSchedule and assign nginx pod to a node

with the taint tz=tanzania:NoSchedule.

SOLU

====

Show current taint :

```
$ k describe no | egrep -i taints:
```

```
Taints:          node-role.kubernetes.io/master:NoSchedule
Taints:          <none>
Taints:          <none>
```

Add a taint to worker node kube3c :

```
$ k taint node kube3c tz=tanzania:NoSchedule
node/kube3c tainted
```

```

$ k describe no | egrep -i taints:
Taints:          node-role.kubernetes.io/master:NoSchedule
Taints:          <none>
Taints:          tz=tanzania:NoSchedule
$

k run nginx --image=nginx --restart=Never --dry-run -o yaml > nginx-pod1.yaml
$
$ cat nginx-pod1.yaml
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx
    name: nginx
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Never
  tolerations:
    <== added this
  - key: "tz"
    operator: "Equal"
    value: "tanzania"
    effect: "NoSchedule"
    <== added this
$
$ k apply -f nginx-pod1.yaml
pod/nginx created

$ k get po -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE      NOMINATED NODE
READINESS GATES
nginx     1/1     Running   0           43s   192.168.80.221  kube3c   <none>
<none>
$

```

AS seen above the pod is deployed on kube3c , which has the taint.

Do the clean up Now..

```

$ k describe no kube3c | egrep -i taint
Taints:          tz=tanzania:NoSchedule
$ k taint node kube3c tz-

```

```

node/kube3c untainted
$ k describe no kube3c | egrep -i taint
Taints:          <none>
$
Deleted the nginx pod

```

```
=====
```

16. Refer to <https://kubernetes.io/docs/tasks/debug-application-cluster/core-metrics-pipeline/>

PRE-REQ : 1. Install Metric Server on the Master Node

2. Run below commands on Master Node :

```

kubect1 run nginx      --image=nginx      --restart=Never -l
app=hello
kubect1 run redis      --image=redis      --restart=Never -l
app=hello
kubect1 run consul     --image=consul     --restart=Never -l
app=hello
kubect1 run memcached  --image=memcached  --restart=Never -l
app=hello

```

Q. From the pods label app: hello , name the pod that shows high CPU and Memory usage ,

dump the output to /tmp/cpu.usage and /tmp/mem.usage

SOLU

```
$ k top po -l app=hello
```

NAME	CPU(cores)	MEMORY(bytes)
consul	14m	23Mi
memcached	1m	5Mi
nginx	0m	2Mi
redis	2m	9Mi

```
$
```

```
$ k top po -l app=hello | sort -nr -k2 | head -1 | awk '{print $1}' | sudo tee /tmp/cpu.usage
```

^^<== refer to 2nd Column

```
$ k top po -l app=hello | sort -nr -k3 | head -1 | awk '{print $1}' | sudo tee /tmp/mem.usage
```


^^<== refer to 3rd Column

I am not sure if they want ONLY the pod name or.....

You can :

\$ k top po consul | tail -1 | awk '{print \$1}' > /tmp/cpu.usage , then edit and only leave 1 pod entry with higher cpu usage

\$ k top po consul | tail -1 | awk '{print \$1}' > /tmp/cpu.memory , then edit and only leave 1 pod entry with higher Mem usage

=====

17. Sort the PV in ascending order from top to bottom.

SOLU

=====

\$ k get pv --sort-by=.spec.capacity.storage

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
all-1	2609192632320m	RWX	Retain	Available	manual
65s					
all-2	2780991324160m	RWX	Retain	Available	manual
65s					
all-3	2791728742400m	RWX	Retain	Available	manual
65s					
all-4	3006477107200m	RWX	Retain	Available	manual
65s					
all-6	6Gi	RWX	Retain	Available	manual
65s					
all-5	6227702579200m	RWX	Retain	Available	manual
65s					

\$

\$

17.a Sort all below pv by name

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
myvolume	10Gi	RWO,RWX	Retain	Available	
normal	6s				
myvolume1	15Gi	RWO,RWX	Retain	Available	
normal	5s				
pv-1f	100Gi	RWO	Retain	Available	
5s					
pv001-1	2Gi	RWO,RWX	Retain	Available	

```
manual                6s
    task-pv-volume    1Gi          RWO,RWX          Retain          Available
manual                6s

SOLU
=====
kube3a k get pv --sort-by=.metadata.name
NAME                CAPACITY    ACCESS MODES    RECLAIM POLICY    STATUS    CLAIM
STORAGECLASS    REASON    AGE
myvolume        10Gi          RWO,RWX          Retain          Available
normal                96s
myvolume1        15Gi          RWO,RWX          Retain          Available
normal                95s
pv001-1          2Gi          RWO,RWX          Retain          Available
manual                96s
pv-1f            100Gi          RWO          Retain          Available
95s
task-pv-volume    1Gi          RWO,RWX          Retain          Available
manual                96s
kube3a
```

=====

18. From etcd pod get the log error msg "store.index" and copy to /opt/KK/xx

```
SOLU
=====
$ k log etcd-kube3a -n kube-system | egrep "store.index" | sudo tee /opt/KK/xx
```

=====

19. Create nginx deployment that has 2 pods, with servicename as nginx-svc ,
exposed thru nodeport.
Resolve Pod IP and Service Name in DNS in any container.

```
SOLU
=====
$ k run nginx --image=nginx --replicas=2 --port=80
$ k expose deployment nginx --type=NodePort --port=80 --target-port=80 --name=nginx-svc
```

After I executed above :

```
$ k get deploy,svc
```

```
NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
deployment.extensions/nginx        2/2      2              2            53s

NAME                                TYPE             CLUSTER-IP      EXTERNAL-IP    PORT(S)          AGE
service/kubernetes                ClusterIP         10.96.0.1       <none>         443/TCP          6d5h
service/nginx-svc                 NodePort          10.103.112.224  <none>         80:30071/TCP     19s
$
```

```
We know the default label for nginx pod is run=nginx , so lets use the default
label and
The best way to get pod IP is :
$ k get po -l run=nginx -o wide
NAME                                READY    STATUS    RESTARTS    AGE    IP              NODE
NOMINATED NODE    READINESS GATES
nginx-57867cc648-7n2tc    1/1      Running    0           11m    192.168.79.226  kube3b
<none>                <none>
nginx-57867cc648-bm2vf    1/1      Running    0           11m    192.168.80.222  kube3c
<none>                <none>
$                                ^^^^^^^^^^^^^^^<== Pod
IP's
```

```
Resolve the pod IP's in consule pod :
$ k exec consul -- nslookup 192.168.79.226
Name:      192.168.79.226
Address 1: 192.168.79.226 192-168-79-226.nginx-svc.default.svc.cluster.local
$
$ k exec consul -- nslookup 192.168.80.222
Name:      192.168.80.222
Address 1: 192.168.80.222 192-168-80-222.nginx-svc.default.svc.cluster.local
$
```

```
Resolve the Service Name , the service name is shown when run k get svc :
$ k exec consul -- nslookup nginx-svc
Name:      nginx-svc
Address 1: 10.103.112.224 nginx-svc.default.svc.cluster.local
$
```

=====
REF : <https://medium.com/@jmarhee/using-initcontainers-to-pre-populate-volume-data-in-kubernetes-99f628cd4519>

```
20. Include busybox initcontainer that creates empty host file in /data , use non
persistent volume.
apiVersion: v1
kind: Pod
```

```

metadata:
  name: my-app
spec:
  containers:
  - name: my-app
    image: nginx
    volumeMounts:
    - name: config-data
      mountPath: /data

```

#

SOLU

=====

I created in steps , in first step I made sure the above part is able to create a pod :

```
$ cat nginx-pod2.yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  volumes:
  - name: config-data
    emptyDir: {}
  containers:
  - name: my-app
    image: nginx
    volumeMounts:
    - name: config-data
      mountPath: /data

```

\$

Above pod worked fine..

Next was to create another busybox pod that would create empty hosts file in /data , where /data is emptyDir

```
$ k run bbox --image=busybox --restart=Never --dry-run -o yaml -- /bin/sh -c 'touch /data/hosts' > bbox11.yaml
```

I extracted info from bbox11.yaml and create my final pod manifest.

```
$ cat nginx-pod2.yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  volumes:

```

```

- name: config-data
  emptyDir: {}
containers:
- name: my-app
  image: nginx
  volumeMounts:
    - name: config-data <== this matches the name under volume:
      mountPath: /data
# below is from busybox container
initContainers:
- args:
  - /bin/sh
  - -c
  - touch /data/hosts
  image: busybox
  name: bbox
  volumeMounts:
    - name: config-data <== this matches the name under volume:
      mountPath: /data
# Above is from busybox container
dnsPolicy: ClusterFirst
restartPolicy: Never

```

After I created above pod :

```
$ k get po
```

NAME	READY	STATUS	RESTARTS	AGE
my-app	1/1	Running	0	10m

```
$ k exec my-app -- ls /data/hosts
/data/hosts
```

```
$ k exec my-app -- ls -l /data/hosts
-rw-r--r-- 1 root root 0 Mar 13 14:30 /data/hosts
```

```
$
Time to delete my pod :
```

```
$ k delete po my-app
pod "my-app" deleted
$
```

REF : <https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/ic/deploy.yaml>

21 . Create alpine init container so can be used with deploy , do create path to

```

local disk.
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: ic-deploy
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: ic
    spec:
      containers:
      - name: main
        image: centos:7
        command:
        - "bin/bash"
        - "-c"
        - "while true; do cat /ic/this; sleep 5; done"
        volumeMounts:
        - mountPath: /ic
          name: msg

```

SOLU

=====

Loking at above manifest, seems like volume is missing , so lets add that , and
run : k run alpine --image=alpine --restart=Never --dry-run -o yaml -- /bin/sh -c
'echo BOSSSS > /ic/this' > alpine.yaml
use info from alpine.yaml file and add in below deployment.

```

$ cat ic-deploy.yaml
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: ic-deploy
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: ic
    spec:
      volumes:
      - name: msg
        hostPath:

```

```

    path: /mnt/data
containers:
- name: main
  image: centos:7
  command:
  - "bin/bash"
  - "-c"
  - "while true; do cat /ic/this; sleep 5; done"
  volumeMounts:
  - mountPath: /ic
    name: msg
initContainers:
- args:
  - /bin/sh
  - -c
  - echo BOSS > /ic/this
  image: alpine
  name: alpine
  volumeMounts:
  - mountPath: /ic
    name: msg

```

```

##      dnsPolicy: ClusterFirst    <== this is part of the alpine pod,
##      restartPolicy: Never      <== was generated when created alpine pod
$

```

```
$ k get deploy,po
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/ic-deploy	1/1	1	1	32m

NAME	READY	STATUS	RESTARTS	AGE
pod/ic-deploy-786f5d879d-pckxg	1/1	Running	0	32m

```
$
```

```
$ k exec ic-deploy-786f5d879d-pckxg -- cat /ic/this
BOSS
```

```
$
```

```
$ k logs ic-deploy-786f5d879d-pckxg
BOSS
```

22. Add initContainer and non persistent disk

```

apiVersion: v1
kind: Pod
metadata:

```

```

name: happypanda
spec:
  containers:
  - name: busybox
    image: busybox:latest
    command: ["/bin/sh", "-c"]
    args: ["cat /opt/workdir/helloworld && sleep 3600"]
    volumeMounts:
    - name: workdir
      mountPath: /opt/workdir

```

SOLU
=====

```

$ k run alpine --image=alpine --restart=Never --dry-run -o yaml -- /bin/sh -c
'echo BOSS > /opt/workdir/helloworld'

```

```

$ cat initcon.yaml
apiVersion: v1
kind: Pod
metadata:
  name: happypanda
spec:
  volumes:
  - name: workdir
    emptyDir: {}    <== this is non persistent disk
#
  containers:
  - name: busybox
    image: busybox
    command: ["/bin/sh", "-c"]
    args: ["cat /opt/workdir/helloworld && sleep 3600"]
    volumeMounts:
    - name: workdir
      mountPath: /opt/workdir
#
# Below was generated from above "k run alpine ....."
initContainers:
- args:
  - /bin/sh
  - -c
  - echo BOSS > /opt/workdir/helloworld
  image: alpine
  name: alpine
  volumeMounts:

```

<== This I copied from above

[illegible]

After I deployed above pod , then :

I ran "k describe po happypanda" , and only showing you few lines..

23. List a Node with high CPU Usage and copy the output to /opt/TTT/cpu.higg.usage
List a Node with High Memory Usage and save the output to
/opt/TTT/mem.usage.high

```

root@kube3a:~#
root@kube3a:~# kubectl top no | sort -rn -k2 | head -1 | awk '{print $1}'
kube3a          ^^<== -k is the column no.

root@kube3a:~# kubectl top no | sort -rn -k2 | head -1 | awk '{print $1}' >
/opt/TTT/cpu.higg.usage
root@kube3a:~# cat /opt/TTT/cpu.higg.usage
kube3a
root@kube3a:~#

```

```

root@kube3a:~# kubectl top no | sort -rn -k4 | head -1 | awk '{print $1}'
kube3a
root@kube3a:~# kubectl top no | sort -rn -k4 | head -1 | awk '{print $1}' >
/opt/TTT/mem.usage.high
root@kube3a:~# cat /opt/TTT/mem.usage.high
kube3a

```

=====

24. Count number of Nodes that are in ready state , except ones which shows tolerance : noschedule

save the results in /opt/TTT/node-ready

SOLU

====

As we don't have perm to write in /opt/TTT dir , for that we'll need to become

root :

\$ cp a /tmp/a

kube3a sudo -i

root@kube3a:~#

root@kube3a:~# source /tmp/a

kube3a

kube3a id -a

uid=0(root) gid=0(root) groups=0(root)

kube3a

copy & paste from top of the Q : kubectl config use-config <cluster_name>

kube3a k get no

The connection to the server localhost:8080 was refused - did you specify the right host or port?

kube3a export KUBECONFIG=/etc/kubernetes/admin.conf

kube3a

kube3a k get no

NAME	STATUS	ROLES	AGE	VERSION
kube3a	Ready	master	14d	v1.13.1
kube3b.home.local	Ready	<none>	14d	v1.13.1
kube3c.home.local	Ready	<none>	9d	v1.13.1

kube3a k describe no | egrep -i taint:

kube3a k describe no | egrep -i taint

Taints: node-role.kubernetes.io/master:NoSchedule

Taints: <none>

Taints: <none>

kube3a

kube3a echo 2 > /opt/TTT/node-ready

kube3a cat /opt/TTT/node-ready

2

kube3a

=====

25. You are unable to Query the cluster nodes :

root@kube3a:~# kubectl get no

The connection to the server localhost:8080 was refused - did you specify the right host or port?

SOLU

=====

If kubelet service is down you can also get above error. Run :

systemctl status -k kubelet

You need to check 2 files , as the /etc/{BROKEN|broken} could be in either of the file.

\$ sudo vi /etc/systemd/system/kubelet.service

OR

\$ sudo vi /lib/systemd/system/kubelet.service

REPLACE : /etc/[BROKEN|broken] OR /etc/kubernetes/BROEKN

WITH : /etc/kubernetes/manifests

\$ sudo systemctl daemon-reload

\$ systemctl status -l kubelet

OR

\$ systemctl restart kubelet

ALso check that the kubelete service is enabled :
\$ systemctl list-unit-files | egrep -i kubelet

Here you'll need to wait for say 2mins , before able to run command kubect1 get no

=====

26. Create a deployment which contains one nginx container. The deployment needs a label
called app:frontend and should be configured for 5 replicas.
Expose a deployment with clusterIP service that listens on tcp/8080 and target tcp/8080.
SOLU
=====

=====

27.
1. create a pod named "web" using image nginx:1.11.9-alpine, on ports 80 and 443
SOLU
=====

```
kube3a k run web --image=nginx:1.11.9-alpine --restart=Never --port=80 --dry-run -o yaml > nG.yaml
```

kube3a

```
kube3a vi nG.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: web
  name: web
spec:
  containers:
  - image: nginx:1.11.9-alpine
    name: web
    ports:
    - containerPort: 80
    - containerPort: 443    <== I manually added this entry
  resources: {}
dnsPolicy: ClusterFirst
```

```
restartPolicy: Never
status: {}
```

```
kube3a k create -f nG.yaml
pod/web created
```

```
kube3a k get po
NAME      READY   STATUS    RESTARTS   AGE
web       1/1     Running   0           15s
kube3a
```

```
kube3a k describe po web | egrep -i port
Ports:          80/TCP, 443/TCP
Host Ports:     0/TCP, 0/TCP
kube3a
```

2. create a service to expose that pod, named as "webservice"

SOLU
=====

```
kube3a k expose pod web --port=80,443 --name=webservice
service/webservice exposed
```

```
kube3a k get svc webservice
NAME            TYPE           CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
webservice      ClusterIP      10.107.250.6    <none>           80/TCP,443/TCP   8s
kube3a
```

3. copy the dns records for the service in file /opt/TTT/web.dnsrecord
record the pod dns in /opt/TTT/pod.dns

SOLU
=====

```
kube3a k run bbox --image=busybox:1.28 --restart=Never -- /bin/sh -c 'sleep 36000'
pod/bbox created
```

```
kube3a k get po -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP             NODE              <none>
NOMINATED NODE   READINESS GATES
bbox      1/1     Running   0           45s   10.244.1.10    kube3b.home.local <none>
<none>
web       1/1     Running   0           11m   10.244.2.25    kube3c.home.local <none>
<none>
kube3a                                     ^^^^^^^^^^<== this is the pod IP , which
```

needs to resolve as pod dns

```
kube3a k exec bbox -- nslookup 10.244.2.25 | sudo tee /opt/TTT/pod.dns
[sudo] password for dharmin:
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        10.244.2.25
Address 1: 10.244.2.25 10-244-2-25.webservice.default.svc.cluster.local
kube3a
```

```
kube3a k get svc
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
kubernetes          ClusterIP           10.96.0.1       <none>           443/TCP          21d
webservice          ClusterIP           10.107.250.6    <none>           80/TCP,443/TCP   12m
kube3a
```

```
Here we resolving DNS name of service webservice :
kube3a k exec  bbox -- nslookup webservice | sudo tee /opt/TTT/web.dnsrecord
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        webservice
Address 1: 10.107.250.6 webservice.default.svc.cluster.local
kube3a
```

=====

One i found which wasn't there was
Finding all nodes which are in ready status (except those where tolerance :
noschedule)
Also - they answer that Sanjay told to restore the faulty node also worked
Sol - open kubelet systemd file n you will see /etc/ BROKEN will be written
Just remove broken n add /Kubernetes/manifests
Do daemon reload n kubelet restart

Note - this will take time
My advice to u will be do write ASAP
All questions are same
I'm sure you will pass
One more thing - we have to only find out ONE node that has HIGH cpu utilization
Not all the nodes

=====

Kubect1 get node OR pod - you'll see Connection Refused Error.

One Master and One Worker Node.

Login to the FIRST Node

Open the Kubelet Service File.

sudo systemctl status kubelet

Open this file : /lib/systemd/system/kubelet.service

Similar to the

-pod-manifest-path

In that file, you'll see

/etc/broken

You've to change is this to /etc/kubernetes/manifest

=====

NOTE :

For metric server you have to include below :

- command:
- /metrics-server
- --kubelet-insecure-tls
- --kubelet-preferred-address-types=InternalIP
- --logtostderr